

Practical Parallelism

Jaimie Murdock

IU Cognitive Science Program
810 Eigenmann Hall
jammurdo@indiana.edu

November 17, 2011

1 Introduction

- Motivation
- Basics
- Synchronization

2 Libraries

- MapReduce
- Hadoop

3 Examples

- OpenMP
- Python

4 Resources

Motivation

Multiprocessing has moved from HPC-only to SMP

Motivation

Multiprocessing has moved from HPC-only to SMP

Cores are cheap:

- 2 – Intel Celeron E3400 – \$47
- 4 – AMD Athlon II X4 631 – \$90
- 6 – AMD Phenom II X6 1035T – \$135
- 8 – AMD FX-8120 – \$210

Motivation

Multiprocessing has moved from HPC-only to SMP

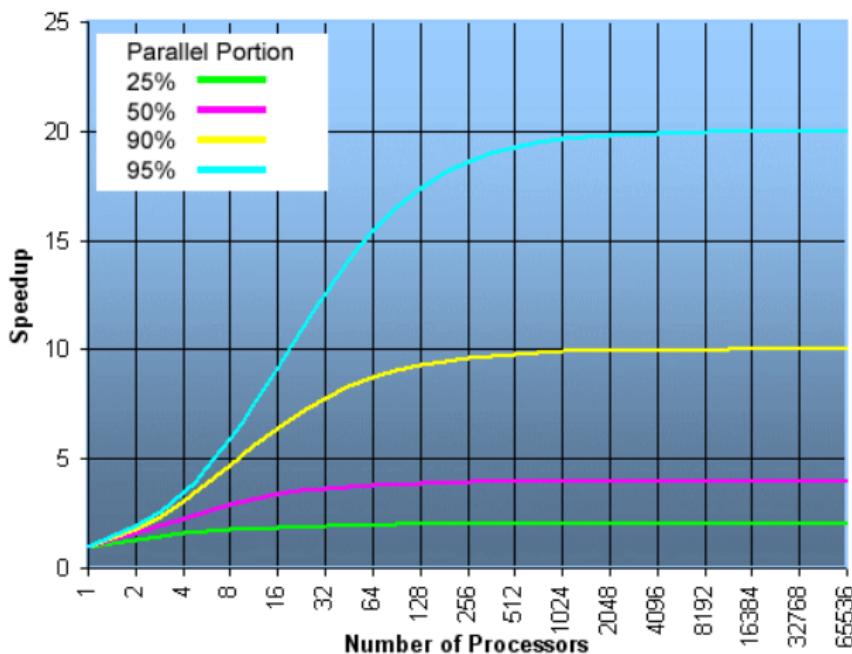
Cores are cheap:

- 2 – Intel Celeron E3400 – \$47
- 4 – AMD Athlon II X4 631 – \$90
- 6 – AMD Phenom II X6 1035T – \$135
- 8 – AMD FX-8120 – \$210

The rise of the cloud — clusters for everyone

How fast?

Amdahl's Law



[fragile]

Painfully Parallel Problems

(or: how to use C211 in an interview)

Painfully Parallel Problems

(or: how to use C211 in an interview)

Given a function that is both commutative and associative (e.g., + or *)

Painfully Parallel Problems

(or: how to use C211 in an interview)

Given a function that is both commutative and associative (e.g., + or *)

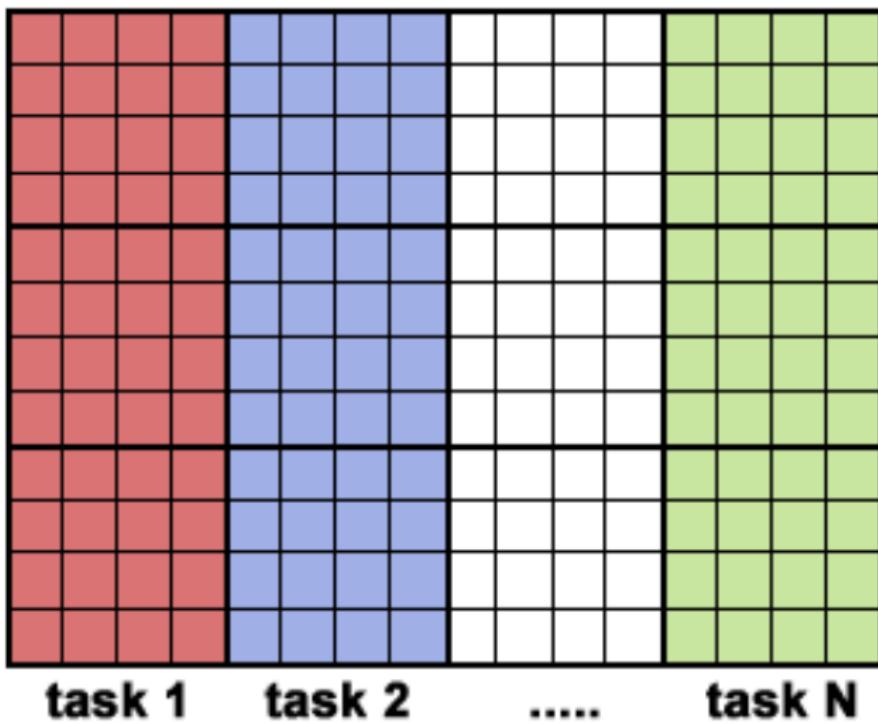
- **Commutative:** $x + y = y + x$
- **Associative:** $(x + y) + z = x + (y + z)$

Partition: 5 6 4 1 4 1 2 5 6 2 7 6 3 4 6 1

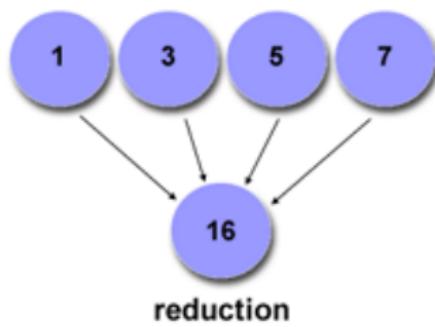
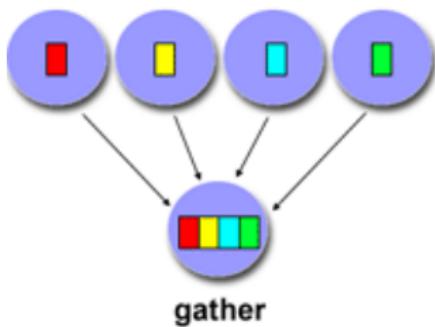
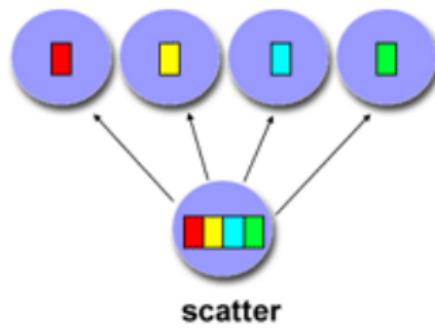
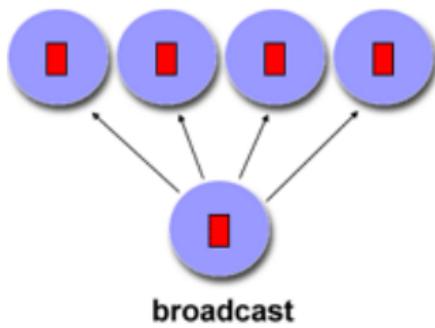
Map: 16 12 21 14

Reduce: 63

Painfully Parallel Problems



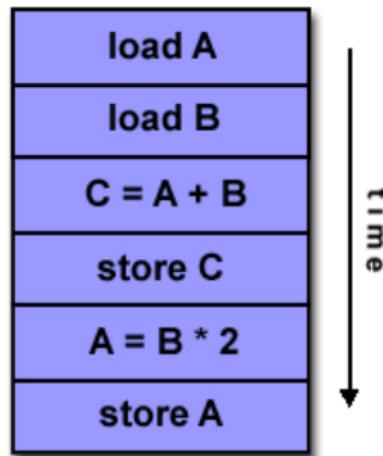
Painfully Parallel Problems



Architecture

Flynn's Taxonomy

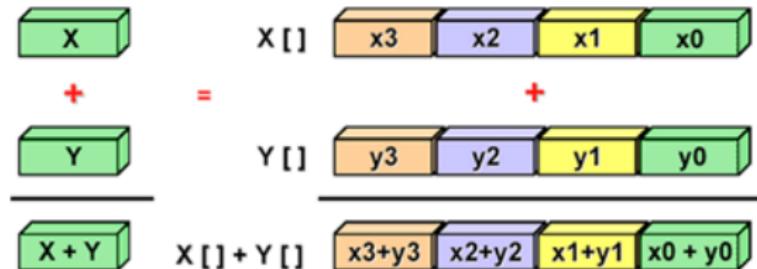
- SISD – Single Instruction, Single Data



Architecture

Flynn's Taxonomy

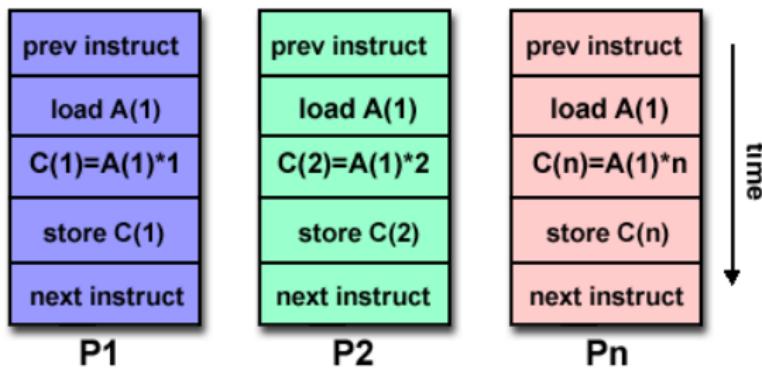
- SISD – Single Instruction, Single Data
- SIMD – Single Instruction, Multiple Data



Architecture

Flynn's Taxonomy

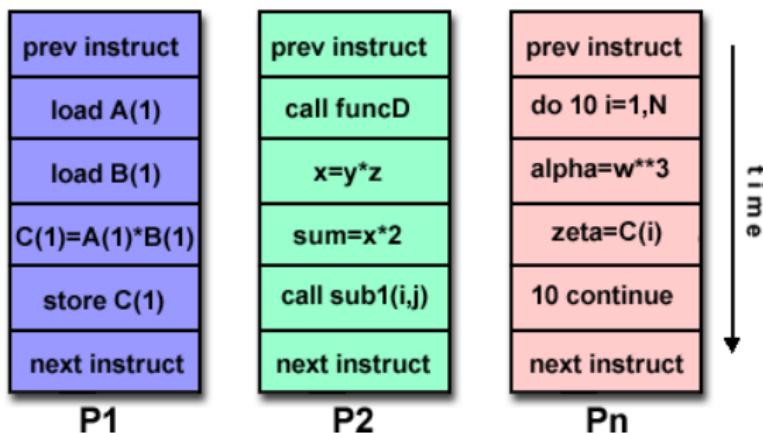
- SISD – Single Instruction, Single Data
- SIMD – Single Instruction, Multiple Data
- MISD – Multiple Instruction, Single Data



Architecture

Flynn's Taxonomy

- SISD – Single Instruction, Single Data
- SIMD – Single Instruction, Multiple Data
- MISD – Multiple Instruction, Single Data
- MIMD – Multiple Instruction, Multiple Data



Synchronization

Process 1

 $a = \text{read}(A)$ $a = a + 1$ $A = \text{write}(a)$

Process 2

 $a = \text{read}(A)$ $b = \text{read}(B)$ $A = \text{write}(a + b)$

Synchronization

Process 1

 $a = \text{read}(A)$ $a = a + 1$ $A = \text{write}(a)$

Process 2

 $a = \text{read}(A)$ $b = \text{read}(B)$ $A = \text{write}(a + b)$

Solutions:

- Barriers
- Locking
- Semaphores

Locks

C#

```
class Account {      // this is a monitor of an account
    long val = 0;

    public void Deposit(const long x) {
        lock (this) {    // only 1 thread at a time may execute this statement
            val += x;
        }
    }

    public void Withdraw(const long x) {
        lock (this) {
            val -= x;
        }
    }
}
```

Locks

Java

```
class Account {      // this is a monitor of an account
    int val = 0;

    public synchronized void Deposit(int x) {
        val += x;
    }

    public synchronized void Withdraw(int x) {
        val -= x;
    }
}
```

Locks

Python

```
from threading import Lock
class Account:
    def __init__(self, value):
        self.value = value
        self.lock = Lock()

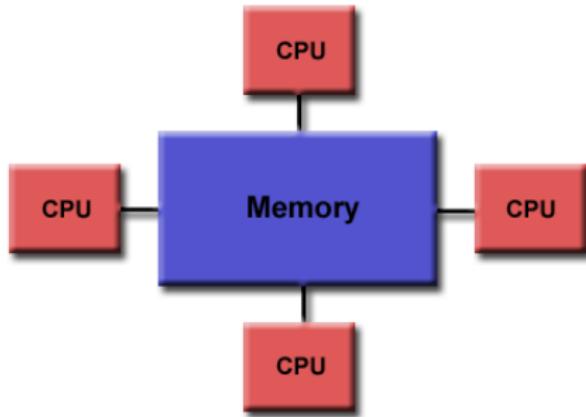
    def deposit(x):
        with self.lock:
            self.value += x

    def withdraw(x):
        with self.lock:
            self.value -= x
```

Memory Models

Shared Memory

- Same address space
- Multithreading



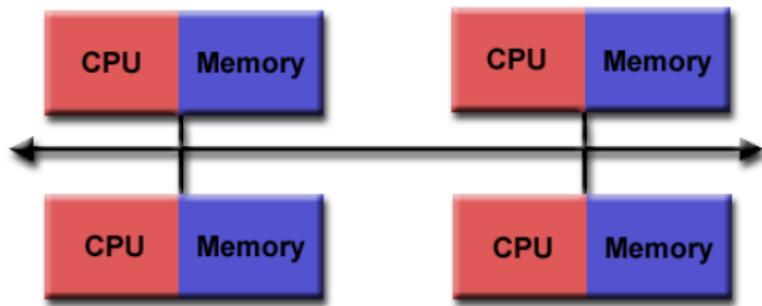
Memory Models

Shared Memory

- Same address space
- Multithreading

Distributed Memory

- Different address space
- Cluster computing/HPC



Memory Models

Shared Memory

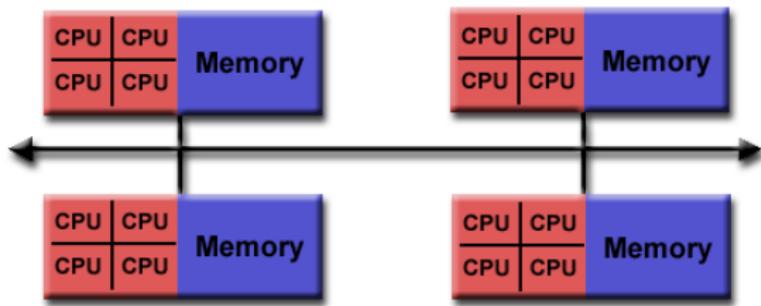
- Same address space
- Multithreading

Distributed Memory

- Different address space
- Cluster computing/HPC

Hybrid Shared-Distributed Memory

- Multiprocessing
- GPU Programming



Threads vs. Processes

Threads

- Shared memory
- Single process (core), multiple execution paths
- Subroutines, GUIs
- Low OS overhead
- Requires synchronization
- POSIX threads, Python `threading`, Java `java.lang.Thread`

Process

- Distributed memory
- Message passing
- High OS overhead
- Only way to utilize multicore or clusters

Pipes and Queues

Queue

- Binary communication of messages
- I/O across processes

Pipes and Queues

Queue

- Binary communication of messages
- I/O across processes

Pipe

- One-way communication of messages.
- Useful for handling I/O to multiple processes

MapReduce

2004 Google framework

<http://labs.google.com/papers/mapreduce.html>

APIs: C++, C#, Erlang, Java, OCaml, Perl, Python, PHP, Ruby, F#, R

MapReduce

2004 Google framework

<http://labs.google.com/papers/mapreduce.html>

APIs: C++, C#, Erlang, Java, OCaml, Perl, Python, PHP, Ruby, F#, R

Six-stage pipeline:

- input reader
- map function
- partition function
- comparison function
- reduce function
- output writer

MapReduce

```
void map(String name, String document):  
    // name: document name  
    // document: document contents  
    for word in document:  
        EmitIntermediate(word, "1");  
  
void reduce(String word, Iterator partialCounts):  
    // word: a word  
    // partialCounts: a list of aggregated partial counts  
    int sum = 0;  
    for pc in partialCounts:  
        sum += ParseInt(pc);  
    Emit(word, AsString(sum));
```

Hadoop

The Big Concepts

- **JVM** – Implemented in Java, commonly used with Clojure and Scala

Hadoop

The Big Concepts

- **JVM** – Implemented in Java, commonly used with Clojure and Scala
- **HDFS** – distributed file system, operates in 64MB chunks.

Hadoop

The Big Concepts

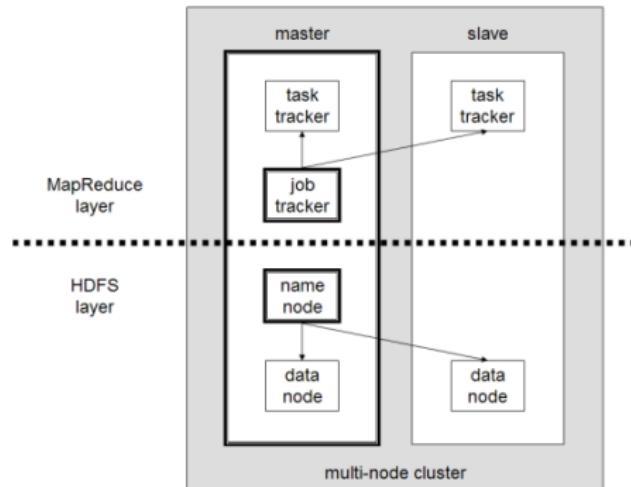
- **JVM** – Implemented in Java, commonly used with Clojure and Scala
- **HDFS** – distributed file system, operates in 64MB chunks.
- **Sharding** – Partitioning. Bring the application to the data to reduce bandwidth

Hadoop

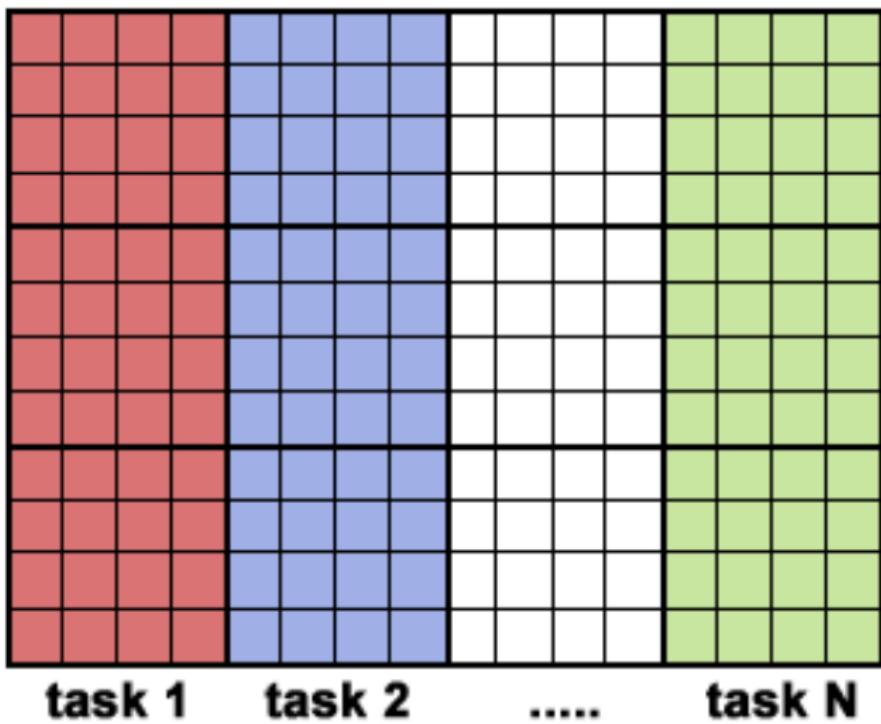
The Big Concepts

- **JVM** – Implemented in Java, commonly used with Clojure and Scala
- **HDFS** – distributed file system, operates in 64MB chunks.
- **Sharding** – Partitioning. Bring the application to the data to reduce bandwidth
- **JobTracker** server takes MapReduce job requests to available **TaskTracker** nodes. (FIFO process pool!)

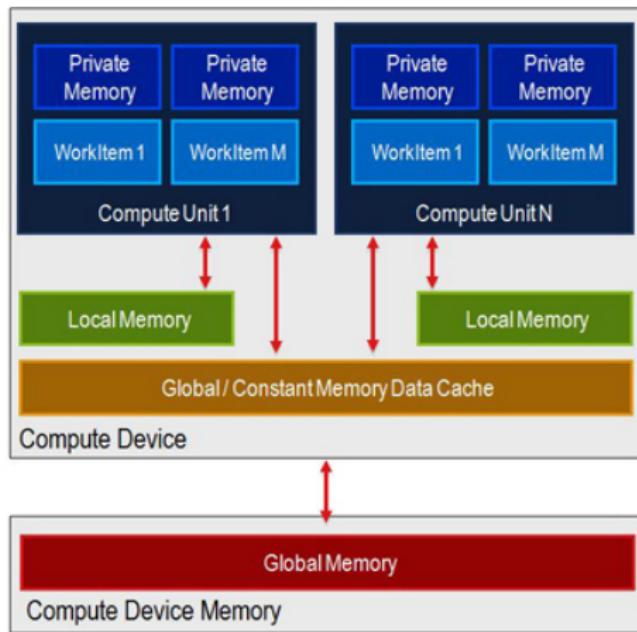
Hadoop



GPU Programming



GPU Memory Model



OpenMP: Upper Triangular Matrix

```
// Calculating Distances
float d; float* D;
for (i = 0; i < POPSIZE; i++) {
    D = dists[i];

    #pragma omp parallel for shared(D, i, dists) private(d, j)
    for (j = i+1; j < POPSIZE; j++) {
        d = distance(i, j);
        D[j] = d;
    }

    D[i] = 0.0;

    #pragma omp parallel for shared(D, i, dists) private(d, j)
    for (j = 0; j < i; j++) {
        d = dists[j][i];
        D[j] = d;
    }
}
```

Python multiprocessing

```
from multiprocessing import Pool  
p = Pool()                      # initialize process pool  
results = p.map(f, args)         # spawn processes  
p.close()                        # close process pool
```

Resources

General

- https://computing.llnl.gov/tutorials/parallel_comp/

Python

- <http://docs.python.org/library/threading.html>
- <http://docs.python.org/library/multiprocessing.html> Python 2.6
- <http://docs.python.org/dev/library/concurrent.futures.html> Python 3.2

Java

- <http://download.oracle.com/javase/7/docs/api/java/lang/Thread.html>
- <http://download.oracle.com/javase/tutorial/essential/concurrency/>

Resources

MapReduce

- <http://labs.google.com/papers/mapreduce.html>
- <http://hadoop.apache.org/>

GPU Programming

- <http://www.khronos.org/opencl/>
- <http://developer.amd.com/zones/openclzone/>
- <http://developer.amd.com/sdks/AMDAPPSDK/samples/>
- <http://developer.nvidia.com/opencl>
- <http://developer.nvidia.com/category/zone/cuda-zone>

1 Introduction

- Motivation
- Basics
- Synchronization

2 Libraries

- MapReduce
- Hadoop

3 Examples

- OpenMP
- Python

4 Resources

Practical Parallelism

Jaimie Murdock

IU Cognitive Science Program
810 Eigenmann Hall
jammurdo@indiana.edu

November 17, 2011